



Initial Prompt

CUDA_KERNEL matmul_atb_kernel

INPUT: A (K×M), B (K×N)

OUTPUT: C (M×N) = $A^T \times B$

GRID: (ceil_div(N,16), ceil_div(M,16)) blocks

BLOCK: 16×16 threads

FOR each thread (global row i = blockIdx.y*16 + ty, global col j = blockIdx.x*16 + tx):



Compile and evaluate this kernel



The performance for this kernel is:

{Speedup}.

Please analyze the provided kernel and try to produce an improved CUDA kernel.



CUDA_KERNEL matmul_atb_optimized:

INPUT: A(K×M), B(K×N)

OUTPUT: C(M×N) = $A^T \times B$

GRID: (M/32, N/32) blocks

BLOCK: 32×32 threads

SHARED_MEM: tile_A[32×33], tile_B[32×33] // +1 padding

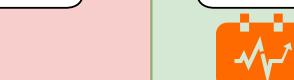
FOR each thread(i,j):

$C[i,j] = \sum_{k=0}^{K-1} LDG(A[k,i]) \times LDG(B[k,j])$

The optimization process for Kevin-32B



Human User



Initial Prompt

CUDA_KERNEL matmul_atb_kernel

INPUT: A (K×M), B (K×N)

OUTPUT: C (M×N) = $A^T \times B$

GRID: (ceil_div(N,16), ceil_div(M,16)) blocks

BLOCK: 16×16 threads

FOR each thread (global row i = blockIdx.y*16 + ty, global col j = blockIdx.x*16 + tx):



Compile and evaluate this kernel



Let's use NCU to see what happened while the kernel was running.



GPU PROFILING RESULTS:

- SM Utilization: 91.94% warp occupancy
- Compute Load: 2.06M active cycles, 444.7M instructions
- Occupancy Bottleneck: Register-limited (4 vs 16 blocks)
- Resource Usage: 53 registers/thread



FIX MEMORY BOTTLENECK (24% stalls):

1. Add shared memory tiling: $(K^3 \times Ci \times \text{float}) \leq 64\text{KB}$
2. Cooperative load + sync + unroll loops
3. Reduce registers to ≤ 32 for 8×128 blocks/SM



Generating the improved kernel following the instruction from Judge...

The process for CudaForge kernel optimization



Coder



Testing



Judge



NCU