

Kevin-32B V.S. CudaForge


Initialization

- Input: {Model Arch in Python}
- Output: {ModelNew}
- {CUDA Strategies}
- {Requirements}



CUDA_KERNEL matmul_atb_kernel
INPUT: A ($K \times M$), B ($K \times N$)
OUTPUT: C ($M \times N$) = $A^T \times B$

GRID: (ceil_div($N, 16$), ceil_div($M, 16$)) blocks
BLOCK: 16×16 threads
FOR each thread (global row $i = blockIdx.y * 16 + ty$, global col $j = blockIdx.x * 16 + tx$):



Human



Testing

Coder



Judge



NCU

Round 1

Kevin-32B

Compile and evaluate this kernel ✓

The performance for this kernel is:
{Speedup}.

Please analyze the provided kernel and
try to produce an improved CUDA kernel.



Compile and evaluate this kernel ✓



Let's use NCU to see what happened while
the kernel was running.



GPU PROFILING RESULTS:

- SM Utilization: 91.94% warp occupancy
- Compute Load: 2.06M active cycles, 444.7M instructions
- Occupancy Bottleneck: Register-limited (4 vs 16 blocks)

Round 2



CUDA_KERNEL matmul_atb_optimized:

INPUT: A($K \times M$), B($K \times N$)

OUTPUT: C($M \times N$) = $A^T \times B$

GRID: ($M/32, N/32$) blocks

BLOCK: 32×32 threads

SHARED_MEM: tile_A[32×33], tile_B[32×33] // +1 padding

FOR each thread(i,j):

$$C[i,j] = \sum_{k=0}^{K-1} LDG(A[k,i]) \times LDG(B[k,j])$$



FIX MEMORY BOTTLENECK (24% stalls):

1. Add shared memory tiling: $(K^3 \times Ci \times float) \leq 64KB$
2. Cooperative load + sync + unroll loops
3. Reduce registers to ≤ 32 for 8×128 blocks/SM



Generating the improved kernel following
the instruction from Judge...